


ASSIGNMENT 5: CONTINUOUS BAG OF WORDS (CBOW) MODEL

The **Continuous Bag of Words (CBOW)** model is a popular word embedding model in deep learning, especially in **Natural Language Processing (NLP)**. It is a part of the **Word2Vec** family, used for learning vector representations of words, which capture semantic and syntactic information. The CBOW model helps understand the relationships between words by predicting a target word based on its surrounding context words.



CBOW Model Explanation

In CBOW, given a set of context words (surrounding words), the model aims to predict the target word. For example, in the sentence: "The cat sat on the mat," if "cat," "on," "the," and "mat" are the context words, the model will try to predict "sat."

Working of the CBOW Model

1. Input Layer:

- The input consists of context words surrounding a target word. For example, if we use a context window of size 2, there would be two words on each side of the target word (4 context words total).
- These context words are represented as **one-hot encoded vectors** in a vocabulary-size dimensional space.

2. Hidden Layer:

- The hidden layer has **no activation function** (it's linear).
- The layer reduces the dimensionality of each word vector to a smaller dimension, which is essentially the size of the embedding vector we want for each word.

3. Projection/Embedding Matrix:

- Each word in the vocabulary is associated with a dense vector (embedding) that is learned by the model.
- The CBOW model uses a shared embedding matrix for all input words.
- When a context word is input, its one-hot vector representation is multiplied by this embedding matrix to get a dense representation.

4. Averaging:

- The embeddings of the context words are **averaged** to get a single context vector, which represents the collective meaning of the surrounding words.

5. Output Layer:

- The output layer is a **softmax** layer that computes probabilities for each word in the vocabulary.
- The word with the highest probability is the model's prediction for the target word.

6. Loss Function:

- CBOW uses the **cross-entropy loss** to calculate the difference between the predicted word probabilities and the actual target word.
- During training, this loss is minimized, updating the embedding vectors to better predict target words based on context.

Important Terms

- **Context Words:** The words surrounding the target word. In CBOW, these words are used to predict the target word.
- **Target Word:** The central word that the model tries to predict.
- **Window Size:** Defines the number of context words considered on either side of the target word. A window size of 2 means two context words before and two after the target word.
- **Word Embedding:** A dense vector representation of a word that captures its meaning based on context. Embeddings are learned during the training process.
- **One-Hot Encoding:** A binary representation of words, where each word is represented by a vector of 0s and 1s, with a unique position set to 1 for each word.
- **Embedding Matrix:** A matrix where each row represents the dense vector of a word in the vocabulary. It is initialized randomly and learned during training.

- **Softmax Function:** Used in the output layer to convert raw scores into probabilities, selecting the target word based on maximum likelihood.
- **Cross-Entropy Loss:** The loss function used to evaluate the difference between the predicted word probability distribution and the actual target word.

This code effectively preprocesses text data, creates a word cloud, generates context-target pairs for CBOW training, and initializes a random embedding matrix. Key terms and steps to note:

- **Stopwords:** Frequently used words removed to focus on meaningful terms.
- **Regular Expressions (Regex):** Used to clean and standardize text.
- **Vocabulary:** The set of unique words in the text.
- **One-Hot Encoding** (conceptual, though not used here): Mapping words to indices.
- **Word Embedding:** Dense vector representation initialized randomly here, but optimized during model training.